

## **Pig: Hadoop Programming Made Easier: Admiring the Pig Architecture, Going with the Pig Latin Application Flow, Working through the ABCs of Pig Latin, Evaluating Local and Distributed Modes of Running Pig Scripts, Checking out the Pig Script Interfaces, Scripting with Pig Latin3.1**

Java MapReduce programs and the Hadoop Distributed File System (HDFS) provide you with a powerful distributed computing framework, but they come with one major drawback — relying on them limits the use of Hadoop to Java programmers who can think in Map and Reduce terms when writing programs.

Pig is a programming tool attempting to have the best of both worlds: a declarative query language inspired by SQL and a low-level procedural programming language that can generate MapReduce code. This lowers the bar when it comes to the level of technical knowledge needed to exploit the power of Hadoop.

Pig was initially developed at Yahoo! in 2006 as part of a research project tasked with coming up with ways for people using Hadoop to focus more on analyzing large data sets rather than spending lots of time writing Java MapReduce programs. The goal here was a familiar one: Allow users to focus more on what they want to do and less on how it's done. Not long after, in 2007, Pig officially became an Apache project. As such, it is included in most Hadoop distributions.

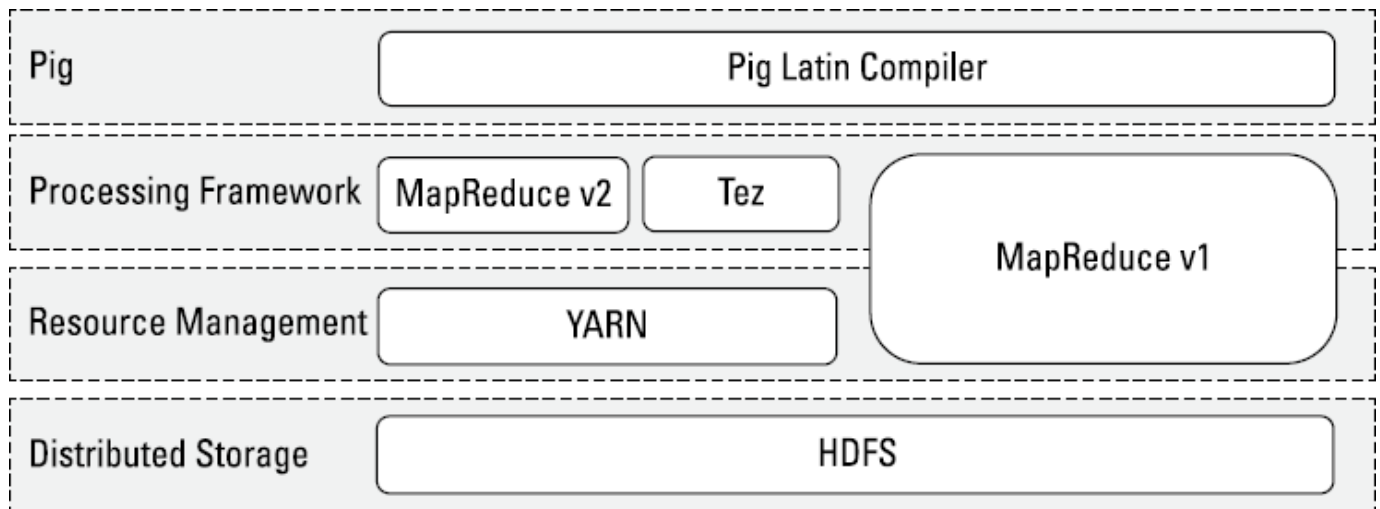
The Pig programming language is designed to handle any kind of data tossed its way — structured, semistructured, unstructured data. Pigs, of course, have a reputation for eating anything they come across. According to the Apache Pig philosophy, pigs eat anything, live anywhere, are domesticated and can fly to boot. Pigs –living anywhere|| refers to the fact that Pig is a parallel data processing programming language and is not committed to any particular parallel framework — including Hadoop. What makes it a domesticated animal? Well, if –domesticated|| means –plays well with humans,|| then it's definitely the case that Pig prides itself on being easy for humans to code and maintain. Lastly, Pig is smart and in data processing lingo this means there is an optimizer that figures out how to do the hard work of figuring out how to get the data quickly. Pig is not just going to be quick — it's going to fly.

### **Admiring the Pig Architecture**

Pig is made up of two components:

- a) The language itself:  
The programming language for Pig is known as Pig Latin, a high-level language that allows you to write data processing and analysis programs.
- b) The Pig Latin compiler:  
The Pig Latin compiler converts the Pig Latin code into executable code. The executable code is either in the form of MapReduce jobs or it can spawn a process where a virtual Hadoop instance is created to run the Pig code on a single node.

The sequence of MapReduce programs enables Pig programs to do data processing and analysis in parallel, leveraging Hadoop MapReduce and HDFS. Running the Pig job in the virtual Hadoop instance is a useful strategy for testing your Pig scripts. Figure 1 shows how Pig relates to the Hadoop ecosystem.



**Figure1: Pig architecture.**

Pig programs can run on MapReduce 1 or MapReduce 2 without any code changes, regardless of what mode your cluster is running. However, Pig scripts can also run using the Tez API instead. Apache Tez provides a more efficient execution framework than MapReduce. YARN enables application frameworks other than MapReduce (like Tez) to run on Hadoop. Hive can also run against the Tez framework.

### **Going with the Pig Latin Application Flow**

At its core, Pig Latin is a dataflow language, where you define a data stream and a series of transformations that are applied to the data as it flows through your application. This is in contrast to a control flow language (like C or Java), where you write a series of instructions. In control flow languages, we use constructs like loops and conditional logic (like an if statement). You won't find loops and if statements in Pig Latin.

If you need some convincing that working with Pig is a significantly easier than having to write Map and Reduce programs, start by taking a look at some real Pig syntax. The following Listing specifies **Sample Pig Code to illustrate the data processing dataflow**.

```
A = LOAD 'data_file.txt';
...
B = GROUP ... ;
...
C= FILTER ...;
...
DUMP B;
..
STORE C INTO 'Results';
```

**Listing: Sample pig code to illustrate the data processing data flow**

Some of the text in this example actually looks like English. Looking at each line in turn, you can see the basic flow of a Pig program. This code can either be part of a script or issued on the interactive shell called Grunt.

- **Load:** You first load (LOAD) the data you want to manipulate. As in a typical MapReduce job, that data is stored in HDFS.

For a Pig program to access the data, you first tell Pig what file or files to use. For that task, you use the LOAD 'data\_file' command. Here, 'data\_file' can specify either an HDFS file or a directory. If a directory is specified, all files in that directory are loaded into the program. If the data is stored in a file format that isn't natively accessible to Pig, you can optionally add the USING function to the LOAD statement to specify a user-defined function that can read in (and interpret) the data.

- **Transform:** You run the data through a set of transformations which are translated into a set of Map and Reduce tasks.

The transformation logic is where all the data manipulation happens.

You can FILTER out rows that aren't of interest, JOIN two sets of data files, GROUP data to build aggregations, ORDER results, and do much, much more.

- **Dump:** Finally, you dump (DUMP) the results to the screen  
or  
Store (STORE) the results in a file somewhere.

### **Working through the ABCs of Pig Latin**

Pig Latin is the language for Pig programs. Pig translates the Pig Latin script into MapReduce jobs that can be executed within Hadoop cluster. When coming up with Pig Latin, the development team followed three key design principles:

- **Keep it simple.**

Pig Latin provides a streamlined method for interacting with Java MapReduce. It's an abstraction, in other words, that simplifies the creation of parallel programs on the Hadoop cluster for data flows and analysis. Complex tasks may require a series of interrelated data transformations — such series are encoded as data flow sequences. Writing data transformation and flows as Pig Latin scripts instead of Java MapReduce programs makes these programs easier to write, understand, and maintain because a) you don't have to write the job in Java, b) you don't have to think in terms of MapReduce, and c) you don't need to come up with custom code to support rich data types.

Pig Latin provides a simpler language to exploit your Hadoop cluster, thus making it easier for more people to leverage the power of Hadoop and become productive sooner.

- **Make it smart.**

You may recall that the Pig Latin Compiler does the work of transforming a Pig Latin program into a series of Java MapReduce jobs. The trick is to make sure that the compiler can optimize the execution of these Java MapReduce jobs automatically, allowing the user to focus on semantics rather than on how to optimize and access the data. SQL is set up as a declarative query that you use to access structured data stored in an RDBMS. The RDBMS engine first translates the query to a data access method and then looks at the statistics and generates a series of data access approaches. The cost-based optimizer chooses the most efficient approach for execution.

- **Don't limit development.**

Make Pig extensible so that developers can add functions to address their particular business problems.

### **Uncovering Pig Latin structures**

The problem we're trying to solve involves calculating the total number of flights flown by every carrier. Following listing is the Pig Latin script we'll use to answer this question.

```
records = LOAD '2013_subset.csv' USING PigStorage(',') AS
    (Year,Month,DayOfMonth,DayOfWeek,DepTime,CRSDep
    Time,ArrTime,CRSArrTime,UniqueCarrier,FlightNum
    ,TailNum,ActualElapsedTime,CRSElapsedTime,AirTi
    me,ArrDelay,DepDelay,Origin,Dest,Distance:int,T
    axiIn,TaxiOut,Cancelled,CancellationCode,Divert
    ed,CarrierDelay,WeatherDelay,NASDelay,SecurityD
    elay,LateAircraftDelay);

milage_recs = GROUP records ALL;
tot_miles = FOREACH milage_recs GENERATE
    SUM(records.Distance);

DUMP tot_miles;
```

Listing : Pig script calculating the total miles flown

The Pig script is a lot smaller than the MapReduce application you'd need to accomplish the same task — the Pig script only has 4 lines of code. And not only is the code shorter, but it's even semi-human readable.

**Most Pig scripts start with the LOAD statement to read data from HDFS.**

In this case, we're loading data from a .csv file. Pig has a data model it uses, so next we need to map the file's data model to the Pig data mode. This is accomplished with the help of the USING statement. We then specify that it is a comma-delimited file with the PigStorage(',') statement followed by the AS statement defining the name of each of the columns.

**Aggregations are commonly used in Pig to summarize data sets.**

The GROUP statement is used to aggregate the records into a single record mileage\_recs. The ALL statement is used to aggregate all tuples into a single group. Note that some statements — including the following SUM statement — requires a preceding GROUP ALL statement for global sums.

**FOREACH... GENERATE statements are used here to transform** columns data In this case, we want to count the miles traveled in the records\_Distance column. The SUM statement computes the sum of the record\_Distance column into a single-column collection total\_miles.

**The DUMP operator is used to execute the Pig Latin statement and display the results on the screen.**

DUMP is used in interactive mode, which means that the statements are executable immediately and the results are not saved. Typically, you will either use the DUMP or STORE operators at the end of your Pig script.

## **Looking at Pig data types and syntax**

Pig's data types make up the data model for how Pig thinks of the structure of the data it is processing. With Pig, the data model gets defined when the data is loaded. Any data you load into Pig from disk is going to have a particular schema and structure. In general terms, though, Pig data types can be broken into two

categories: scalar types and complex types. Scalar types contain a single value, whereas complex types contain other types, such as the Tuple, Bag, and Map types.

Pig Latin has these four types in its data model:

**Atom:** An atom is any single value, such as a string or a number — `__Diego'`, for example. Pig's atomic values are scalar types that appear in most programming languages — `int`, `long`, `float`, `double`, `chararray`, and `bytearray`, for example. See Figure 2 to see sample atom types.

**Tuple:** A tuple is a record that consists of a sequence of fields. Each field can be of any type — `__Diego'`, `__Gomez'`, or `6`, for example. Think of a tuple as a row in a table.

**Bag:** A bag is a collection of non-unique tuples. The schema of the bag is flexible — each tuple in the collection can contain an arbitrary number of fields, and each field can be of any type.

**Map:** A map is a collection of key value pairs. Any type can be stored in the value, and the key needs to be unique. The key of a map must be a `chararray` and the value can be of any type.

Figure -2 offers some fine examples of Tuple, Bag, and Map data types, as well.

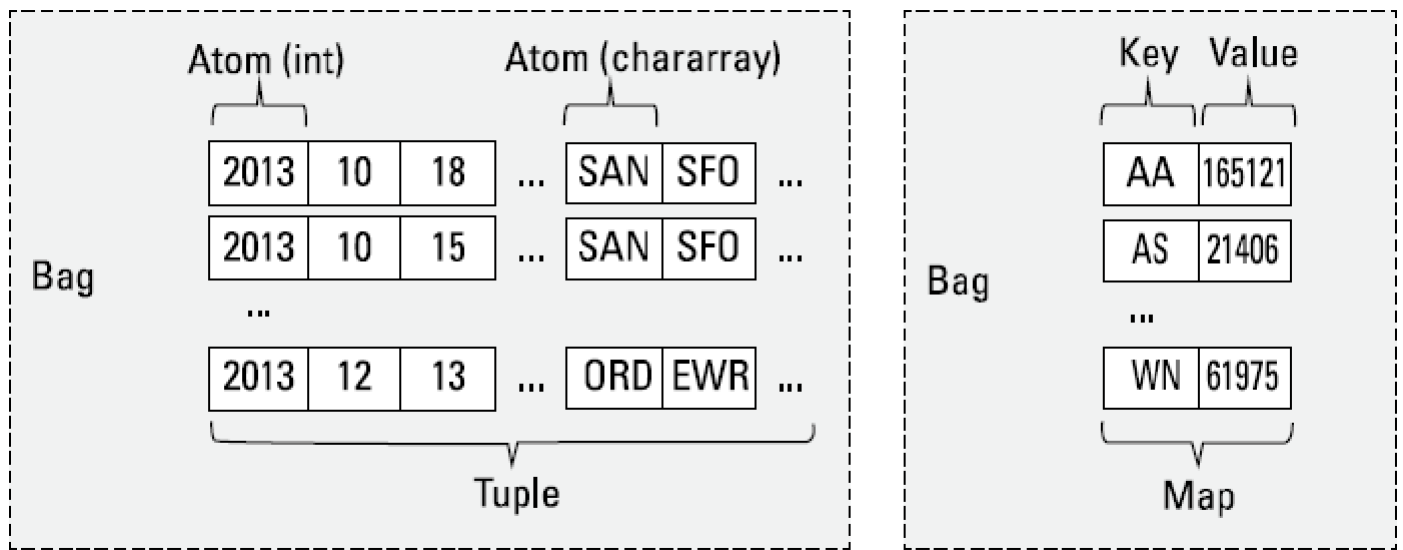


Figure 2: Sample Pig Data Types

In a Hadoop context, accessing data means allowing developers to load, store, and stream data, whereas transforming data means taking advantage of Pig's ability to group, join, combine, split, filter, and sort data. Table 1 gives an overview of the operators associated with each operation.

<i><b>Operation</b></i>	<i><b>Operator</b></i>	<i><b>Explanation</b></i>
Data Access	LOAD/STORE	Read and Write data to file system
	DUMP	Write output to standard output (stdout)
	STREAM	Send all records through external binary
	FOREACH	Apply expression to each record and output one or more records
	FILTER	Apply predicate and remove records that don't meet condition
	GROUP/ COGROUP	Aggregate records with the same key from one or more inputs
	JOIN	Join two or more records based on a condition
Transformations	CROSS	Cartesian product of two or more inputs
	ORDER	Sort records based on key
	DISTINCT	Remove duplicate records
	UNION	Merge two data sets
	SPLIT	Divide data into two or more bags based on predicate
	LIMIT	subset the number of records

Table 1 Pig Latin Operators

Pig also provides a few operators that are helpful for debugging and troubleshooting, as shown in Table 2:

<i><b>Operation</b></i>	<i><b>Operator</b></i>	<i><b>Description</b></i>
Debug	DESCRIBE	Return the schema of a relation.
	DUMP	Dump the contents of a relation to the screen.
	EXPLAIN	Display the MapReduce execution plans.

Table 2 Operators for Debugging and Troubleshooting

The optional USING statement defines how to map the data structure within the file to the Pig data model — in this case, the PigStorage () data structure, which parses delimited text files. The optional AS clause defines a schema for the data that is being mapped. If you don't use an AS clause, you're basically telling the default LOAD Func to expect a plain text file that is tab delimited.



### **Evaluating Local and Distributed Modes of Running Pig scripts**

Before you can run your first Pig script, you need to have a handle on how Pig programs can be packaged with the Pig server. Pig has two modes for running scripts, as shown in Figure 3.

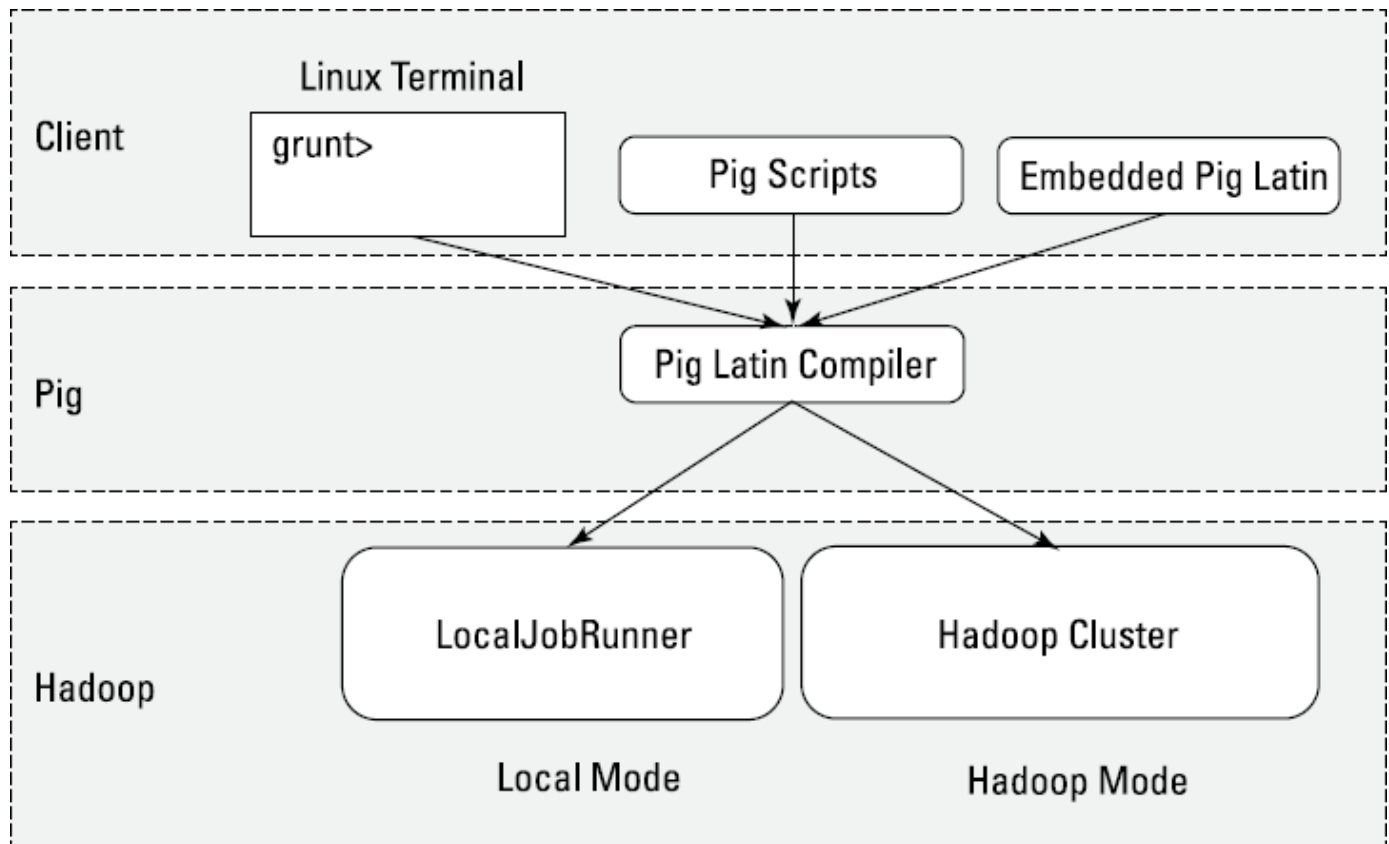


Figure 3. Pig modes

#### **Local mode**

All scripts are run on a single machine without requiring Hadoop MapReduce and HDFS. This can be useful for developing and testing Pig logic. If you're using a small set of data to develop or test your code, then local mode could be faster than going through the MapReduce infrastructure.

Local mode doesn't require Hadoop. When you run in Local mode, the Pig program runs in the context of a local Java Virtual Machine, and data access is via the local file system of a single machine. Local mode is actually a local simulation of MapReduce in Hadoop's LocalJobRunner class.

#### **MapReduce mode (also known as Hadoop mode)**

Pig is executed on the Hadoop cluster. In this case, the Pig script gets converted into a series of MapReduce jobs that are then run on the Hadoop cluster. If you have a terabyte of data that you want to perform operations on and you want to interactively develop a program, you may soon find things slowing down considerably, and you may start growing your storage.

Local mode allows you to work with a subset of your data in a more interactive manner so that you can figure out the logic (and work out the bugs) of your Pig program. After you have things set up as you want them and your operations are running smoothly, you can then run the script against the full data set using MapReduce mode.

### **Checking Out the Pig Script Interfaces**

Pig programs can be packaged in three different ways:

- **Script:** This method is nothing more than a file containing Pig Latin commands, identified by the .pig suffix (FlightData.pig, for example). Ending your Pig program with the .pig extension is a convention but not required. The commands are interpreted by the Pig Latin compiler and executed in the order determined by the Pig optimizer.
- **Grunt:** Grunt acts as a command interpreter where you can interactively enter Pig Latin at the Grunt command line and immediately see the response. This method is helpful for prototyping during initial development and with what-if scenarios.
- **Embedded:** Pig Latin statements can be executed within Java, Python, or JavaScript programs.

Pig scripts, Grunt shell Pig commands, and embedded Pig programs can run in either Local mode or MapReduce mode. The Grunt shell provides an interactive shell to submit Pig commands or run Pig scripts. To start the Grunt shell in Interactive mode, just submit the command pig at your shell. To specify whether a script or Grunt shell is executed locally or in Hadoop mode just specify it in the -x flag to the pig command. The following is an example of how you'd specify running your Pig script in local mode:

```
pig -x local milesPerCarrier.pig
```

Here's how you'd run the Pig script in Hadoop mode, which is the default if you don't specify the flag:

```
pig -x mapreduce milesPerCarrier.pig
```

By default, when you specify the pig command without any parameters, it starts the Grunt shell in Hadoop mode. If you want to start the Grunt shell in local mode just add the -x local flag to the command.

Here is an example:

```
pig -x local
```

### **Scripting with Pig Latin**

Hadoop is a rich and quickly evolving ecosystem with a growing set of new applications. Rather than try to keep up with all the requirements for new capabilities, Pig is designed to be extensible via user-defined functions, also known as UDFs. UDFs can be written in a number of programming languages, including Java, Python, and JavaScript. Developers are also posting and sharing a growing collection of UDFs online. (Look for Piggy Bank and DataFu, to name just two examples of such online collections.) Some of the Pig UDFs that are part of these repositories are LOAD/STORE functions (XML, for example), date time functions, text, math, and stats functions.

Pig can also be embedded in host languages such as Java, Python, and JavaScript, which allows you to integrate Pig with your existing applications. It also helps overcome limitations in the Pig language. One of the most commonly referenced limitations is that Pig doesn't support control flow statements: if/else, while loop, for loop, and condition statements. Pig natively supports data flow, but needs to be embedded within another language to provide control flow. There are tradeoffs, however of embedding Pig in a control-flow language. For example if a Pig statement is embedded in a loop, every time the loop iterates and runs the Pig statement, this causes a separate MapReduce job to run.